

# Machine Learning and Heterogeneous Effects

*MIXTAPE SESSION*

Prof. Brigham Frandsen



# Allow me to introduce myself

- ▶ Economics professor at Brigham Young University in Utah



## Allow me to introduce myself

- ▶ Economics professor at Brigham Young University in Utah
- ▶ 4 biological kids, 3 foster daughters, most of whom can now run and mountain bike faster than me



# Allow me to introduce myself

- ▶ Economics professor at Brigham Young University in Utah
- ▶ 4 biological kids, 3 foster daughters, most of whom can now run and mountain bike faster than me
- ▶ A big fan of causal inference in observational settings:
  - ▶ Quasi-experimental evaluations of the effects of unions  
(Frandsen 2016, 2017, 2021; Chen, Frandsen, Grabowski, Town, Sojourner 2015)
  - ▶ Distributional effects  
(Frandsen and Lefgren 2018, 2021; Frandsen, Froelich, Melly 2012)
- ▶ And of exploring machine learning in applied economics:
  - ▶ Teach Machine Learning for Economists at BYU
  - ▶ Research on the power of ML in empirical strategies  
(Angrist and Frandsen 2022)

# Effects Ex Machina: Where we're going

## **Machine Learning + Heterogeneous Treatment Effects**

- ▶ Causality primer/review
- ▶ Machine learning (ML) prediction primer/review
- ▶ Heterogeneous treatment effects
  - ▶ When they matter
  - ▶ Conceptual framework
  - ▶ Using ML to predict treatment effects:  
Random Causal Forests
  - ▶ Python/R implementation

(Prequel to this course: Machine Learning and Causal Inference)

# Potential outcomes and treatment effects



# Potential outcomes and treatment effects



\$\$\$\$

\$43M

# Potential outcomes and treatment effects



\$\$\$\$

\$43M

## Potential outcomes and treatment effects



\$\$\$\$

\$43M



\$

\$700K

# Potential outcomes and treatment effects



\$\$\$

\$43M

D = 1


Y(1)

\$

\$700K

D = 0

Y(0)



## Potential outcomes and treatment effects



\$\$\$

\$

$$\$43M - \$700K = \$42.3M$$

$$Y(1) - Y(0) = \text{Treatment effect}$$

# Potential outcomes and treatment effects



\$\$\$\$

\$43M

\$

\$700K = \$42.3M

Y(1)

-

Y(0)

= Treatment effect

counterfactual

# Potential outcomes and treatment effects



\$\$\$\$

\$43M

\$

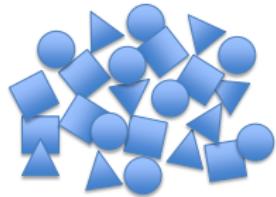
\$700K = \$42.3M

Y(1)


Y(0)

= Treatment effect


# Potential outcomes and treatment effects




# Potential outcomes and treatment effects



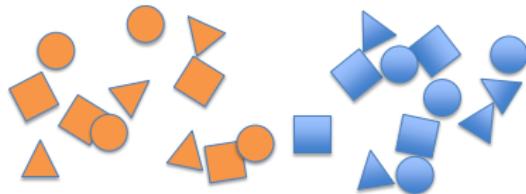
# Potential outcomes and treatment effects



# Potential outcomes and treatment effects




# Potential outcomes and treatment effects

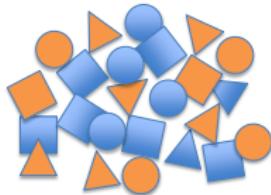


$$E[Y(1)] - E[Y(0)] = \text{ATE}$$

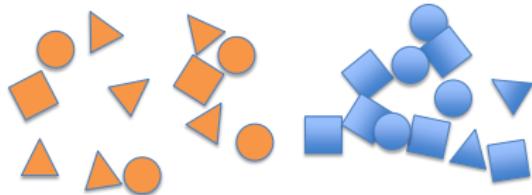
# Potential outcomes and treatment effects



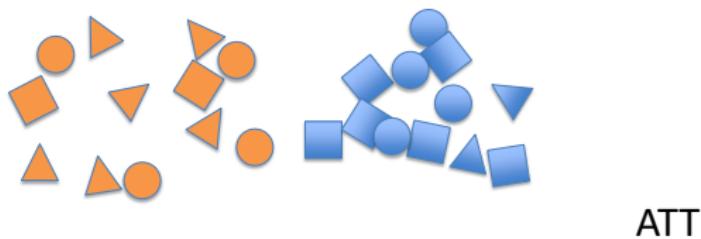
# Potential outcomes and treatment effects



# Potential outcomes and treatment effects



# Potential outcomes and treatment effects

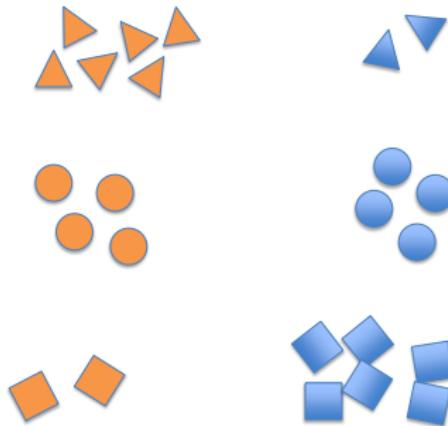


$$E[Y|D=1] - E[Y|D=0] = E[Y(1) - Y(0)|D=1]$$

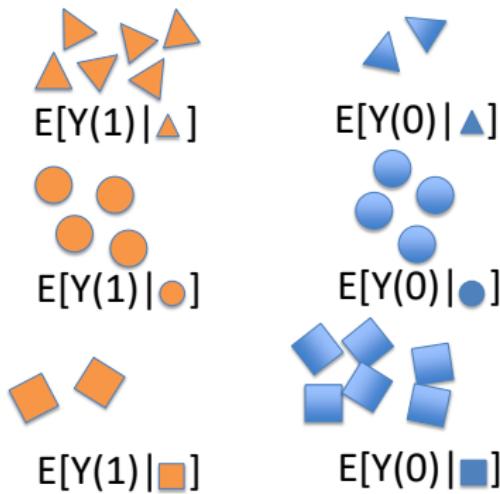
$$+ E[Y(0)|D=1] - E[Y(0)|D=0]$$

selection bias


# Potential outcomes and treatment effects





# Potential outcomes and treatment effects



# Potential outcomes and treatment effects



# Potential outcomes and treatment effects



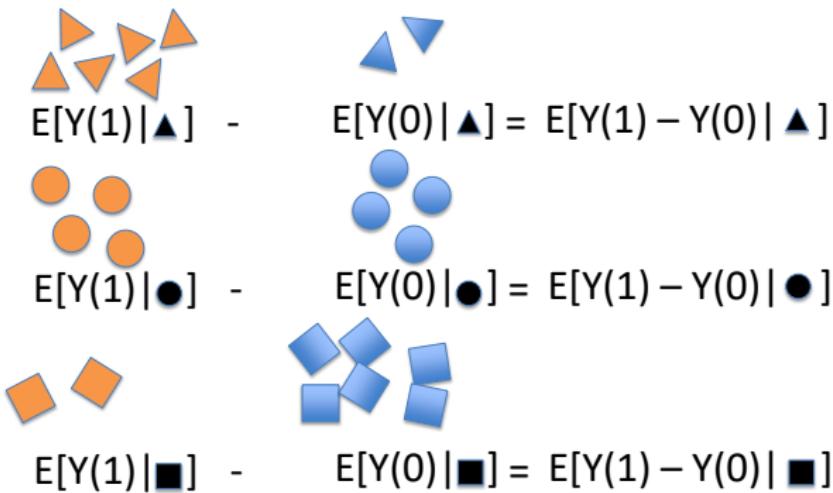
As long as:

$$E[Y(1)|\Delta] = E[Y(1)|\Delta]$$
$$E[Y(0)|\Delta] = E[Y(0)|\Delta]$$


$$E[Y(1)|\bullet] = E[Y(1)|\bullet]$$
$$E[Y(0)|\bullet] = E[Y(0)|\bullet]$$

$$E[Y(1)|\blacksquare] = E[Y(1)|\blacksquare]$$
$$E[Y(0)|\blacksquare] = E[Y(0)|\blacksquare]$$

# Potential outcomes and treatment effects



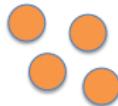
# Potential outcomes and treatment effects



# Potential outcomes and treatment effects



$Y(1), Y(0) \perp\!\!\!\perp D | \blacktriangle$



$Y(1), Y(0) \perp\!\!\!\perp D | \bullet$



$Y(1), Y(0) \perp\!\!\!\perp D | \blacksquare$

# Potential outcomes and treatment effects


$$Y(1), Y(0) \perp\!\!\!\perp D | \blacktriangle$$

$$Y(1), Y(0) \perp\!\!\!\perp D | \bullet$$

$$Y(1), Y(0) \perp\!\!\!\perp D | \blacksquare$$

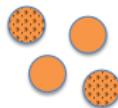
Selection on observables,  
“unconfoundedness”:

$$Y(1), Y(0) \perp\!\!\!\perp D | X$$

# Potential outcomes and treatment effects



$Y(1), Y(0) \perp\!\!\!\perp D | \Delta$




$Y(1), Y(0) \perp\!\!\!\perp D | \bullet$

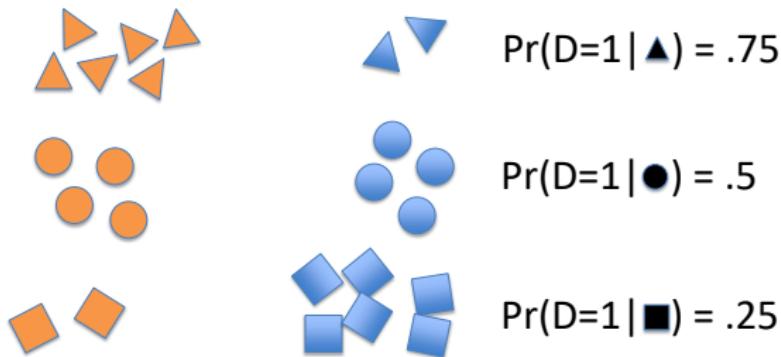


$Y(1), Y(0) \perp\!\!\!\perp D | \blacksquare$

## Potential outcomes and treatment effects



## Potential outcomes and treatment effects

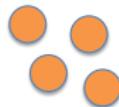


Common support,  
“overlap”:  $0 < \Pr(D=1 | X) < 1$

## Potential outcomes and treatment effects



$$\Pr(D=1 | \blacktriangle) = .75$$



$$\Pr(D=1 | \bullet) = .5$$




$$\Pr(D=1 | \blacksquare) = .25$$



# Potential outcomes and treatment effects



$$\Pr(D=1 | \blacktriangle) = .75$$



$$\Pr(D=1 | \bullet) = .5$$



$$\Pr(D=1 | \blacksquare) = .25$$



!!

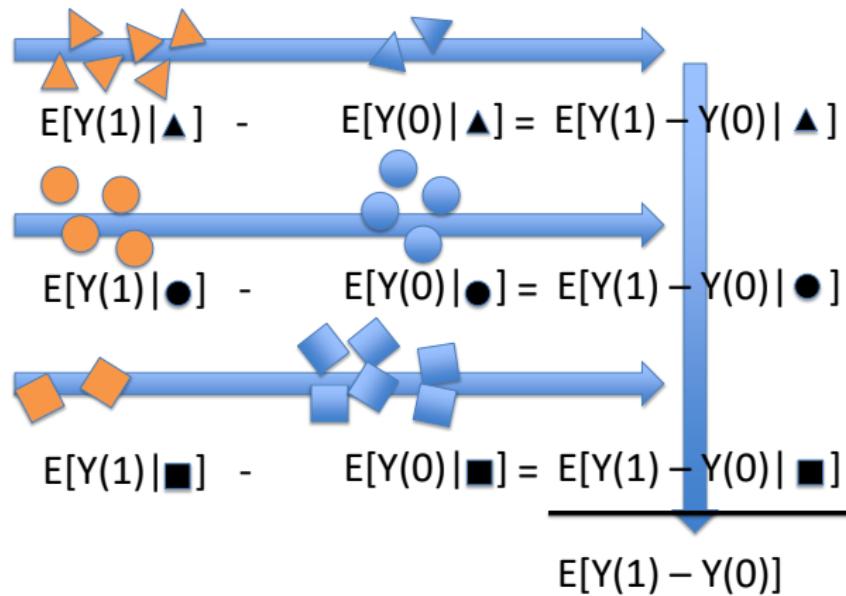
$$\Pr(D=1 | \lozenge) = 1$$

!!



$$\Pr(D=1 | \blacksquare) = 0$$

## Potential outcomes and treatment effects



## Basic causal inference summary

- ▶ Target (for now!):

$$ATE = E [Y_i(1) - Y_i(0)] = E [\tau_i]$$

- ▶ Key identifying assumption:

$$(Y_i(0), Y_i(1)) \perp\!\!\!\perp D_i | X_i$$

- ▶ Estimation:

- ▶ Multiple linear regression (OLS)


$$Y_i = \beta_0 + \tau D_i + \beta_1 X_{1i} + \cdots + \beta_k X_{ki} + \varepsilon$$

- ▶ Matching
  - ▶ Propensity score methods
  - ▶ Machine-assisted:
    - ▶ Post-Double Selection Lasso
    - ▶ Double/De-biased Machine Learning

- ▶ Go to python!


## Prediction Target

$$y_i = \alpha + \beta x_i + \varepsilon_i$$



# Prediction Target

$$y_i = \underbrace{\alpha + \beta x_i}_{\hat{y}} + \varepsilon_i$$




# Prediction Methods

Supervised machine learning algorithms:

- ▶ Decision trees
- ▶ Random forests
- ▶ Penalized regression (ridge, lasso)
- ▶ Support vector machines

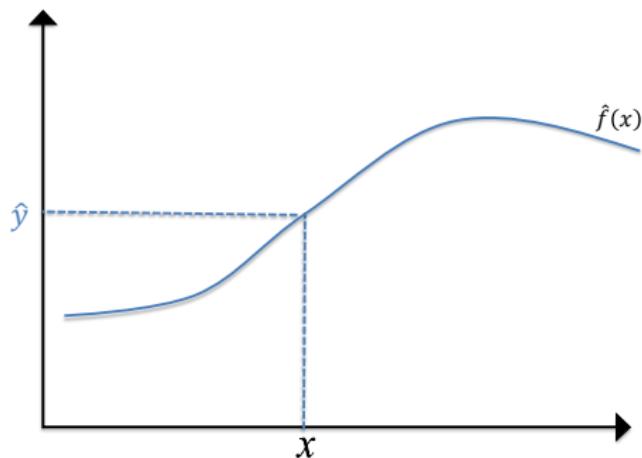
# Prediction mechanics

- ▶ **Goal:** Predict an out-of-sample outcome  $Y$



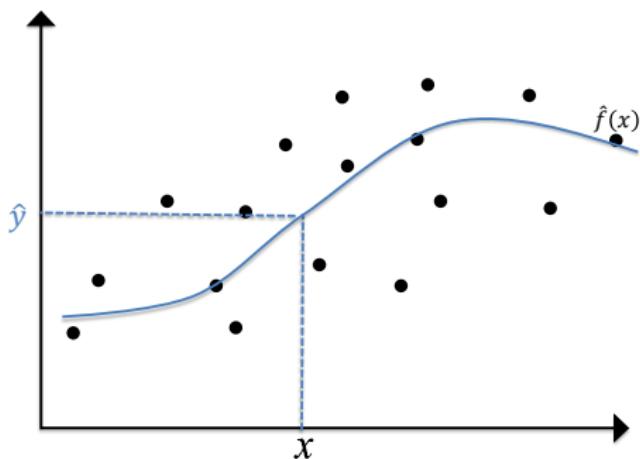
# Prediction mechanics

- ▶ **Goal:** Predict an out-of-sample outcome  $Y$
- ▶ as a function,  $\hat{f}(X)$ , of **features**  $X = (1, X_1, X_2, \dots, X_K)'$ .



# Prediction mechanics


- ▶ **Goal:** Predict an out-of-sample outcome  $Y$
- ▶ as a function,  $\hat{f}(X)$ , of **features**  $X = (1, X_1, X_2, \dots, X_K)'$ .
- ▶ Estimate the function  $\hat{f}$  (aka “train the model”) based on **training sample**  $\{(Y_i, X_i); i = 1, \dots, N\}$



# What's a “good” prediction?

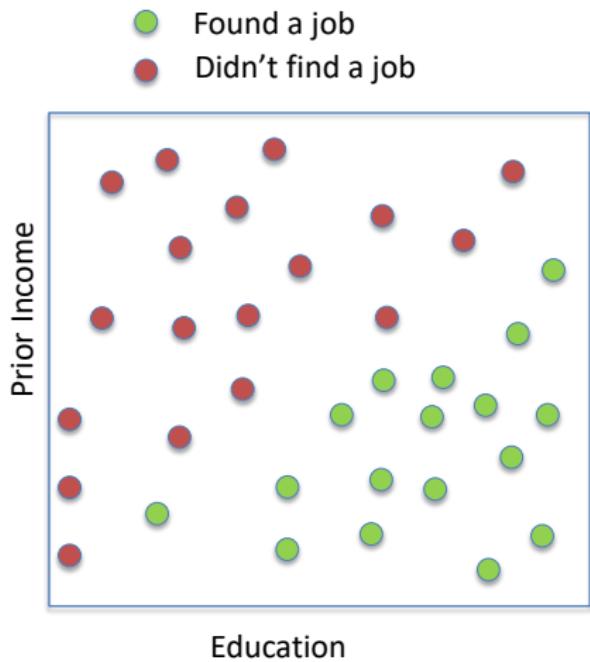
- Want our prediction to be “close,” i.e. minimize the expected **mean squared error**:

$$\min_{f(x)} E \left[ (f(x) - Y)^2 \middle| X = x \right]$$





# Decision Trees

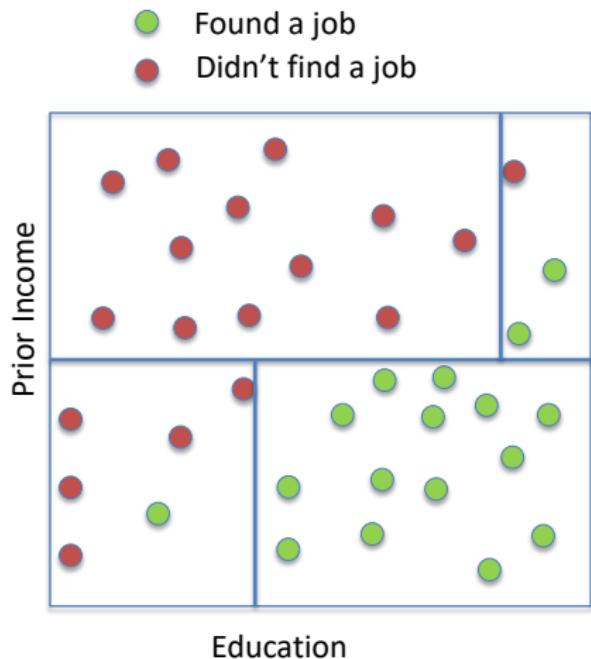

Initial node



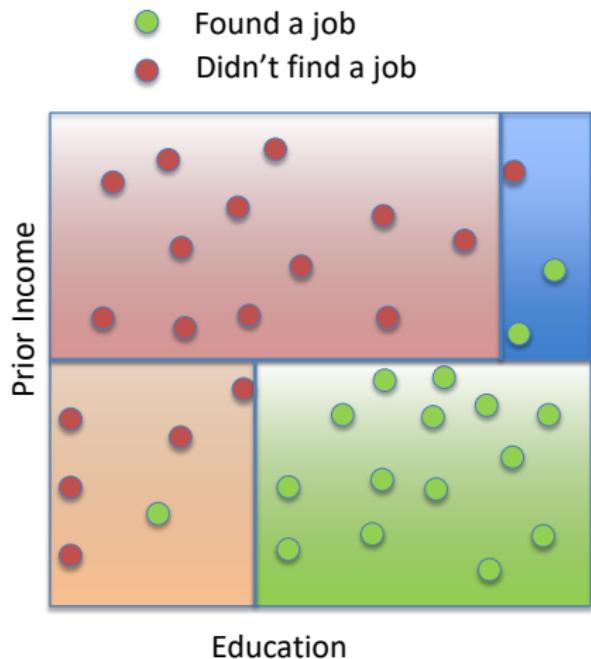

# Decision Trees




# Decision Trees



# Decision Trees




# Decision Trees



- ▶ Where to split:  
Choose the feature from  $\{x_1, \dots, x_p\}$  and the value of that feature to minimize MSE in the resulting child nodes
- ▶ Tuning parameters
  - ▶ Max depth
  - ▶ Min training obs per leaf
  - ▶ Min improvement in fit in order to go ahead with a split

# Forest for the Trees



- ▶ Value proposition: reduce variance by averaging together multiple predictions
- ▶ The catch: individual trees need to be **de-correlated**
- ▶ Algorithm:
  - ▶ Grow  $B$  trees, each on a different bootstrapped sample
  - ▶ At each split, consider only a random subset of features
  - ▶ Average together the individual predictions
- ▶ Let's grow some trees in python!

# Combining causal effects and ML: predicting heterogeneous treatment effects

- ▶ What is the effect of job training on the probability of finding a job . . .
  - ▶ for more-educated vs. less-educated individuals?
  - ▶ for men vs. women?
  - ▶ for married vs. single?
  - ▶ for high-earning vs. low-earning (prior to training)?
  - ▶ for minorities vs. non-minorities?
- ▶ Why does it matter?
- ▶ Other examples where heterogeneity in treatment effects matter?

## Traditional heterogeneity analysis: Interacted regression

To estimate the overall average effect:

$$Y_i = \tau D_i + \varepsilon_i, \quad i \in \{1, \dots, n\}$$

## Traditional heterogeneity analysis: Interacted regression

To estimate the overall average effect:

$$Y_i = \tau D_i + \varepsilon_i, \quad i \in \{1, \dots, n\}$$

To explore heterogeneity by sex:

$$\begin{aligned} Y_i &= \tau^{\text{female}} D_i + \varepsilon_i, & i : \text{Female}_i = 1 \\ Y_i &= \tau^{\text{male}} D_i + \varepsilon_i, & i : \text{Female}_i = 0, \end{aligned}$$

## Traditional heterogeneity analysis: Interacted regression

To estimate the overall average effect:

$$Y_i = \tau D_i + \varepsilon_i, \quad i \in \{1, \dots, n\}$$

To explore heterogeneity by sex:

$$Y_i = \tau^{\text{female}} D_i + \varepsilon_i, \quad i : \text{Female}_i = 1$$

$$Y_i = \tau^{\text{male}} D_i + \varepsilon_i, \quad i : \text{Female}_i = 0,$$

or, equivalently:

$$\begin{aligned} Y_i &= \tau^{\text{male}} D_i + \beta \text{Female}_i + \gamma D_i \times \text{Female}_i + \varepsilon_i; \\ \tau^{\text{female}} &= \tau^{\text{male}} + \gamma. \end{aligned}$$

## Traditional heterogeneity analysis: Interacted regression

To estimate the overall average effect:

$$Y_i = \tau D_i + \varepsilon_i, \quad i \in \{1, \dots, n\}$$

To explore heterogeneity by sex:

$$Y_i = \tau^{\text{female}} D_i + \varepsilon_i, \quad i : \text{Female}_i = 1$$

$$Y_i = \tau^{\text{male}} D_i + \varepsilon_i, \quad i : \text{Female}_i = 0,$$

or, equivalently:

$$\begin{aligned} Y_i &= \tau^{\text{male}} D_i + \beta \text{Female}_i + \gamma D_i \times \text{Female}_i + \varepsilon_i; \\ \tau^{\text{female}} &= \tau^{\text{male}} + \gamma. \end{aligned}$$

More generally,

$$Y_i = \tau D_i + X'_i \beta + D_i X'_i \gamma + \varepsilon_i,$$

$$\tau(x) = \tau + x' \gamma$$

# Challenges with traditional heterogeneity analysis

$$Y_i = \tau D_i + X'_i \beta + D_i X'_i \gamma + \varepsilon_i$$

- ▶ Functional form: treatment effects may not vary linearly with  $X_i$
- ▶ Curse of dimensionality: when  $X_i$  includes many variables, OLS impractical or infeasible
- ▶ These are problems ML was born to solve!

# Predicting outcomes vs. treatment effects

Predicting outcomes

Predicting treatment effects

---

Target:  $\hat{y}(x) = E[Y_i|X_i = x]$

Target:  $\tau(x) = E[\tau_i|X_i = x]$

Criterion:

Criterion:

$$\min E \left[ (\hat{y}(x) - Y_i)^2 | X_i = x \right]$$

$$\min E \left[ (\tau(x) - \tau_i)^2 | X_i = x \right]$$

Training data:  $\{Y_i, X_i\}_{i=1}^n$

Training data:  $\{\tau_i, X_i\}_{i=1}^n$

# Predicting outcomes vs. treatment effects

Predicting outcomes

Predicting treatment effects

---

Target:  $\hat{y}(x) = E[Y_i | X_i = x]$

Target:  $\tau(x) = E[\tau_i | X_i = x]$

Criterion:

Criterion:

$$\min E \left[ (\hat{y}(x) - Y_i)^2 | X_i = x \right]$$

$$\min E \left[ (\tau(x) - \tau_i)^2 | X_i = x \right]$$

Training data:  $\{Y_i, X_i\}_{i=1}^n$

Training data:  $\{\tau_i, X_i\}_{i=1}^n$

Why is training data a problem for predicting treatment effects?

## Predicting outcomes vs. treatment effects

Predicting outcomes	Predicting treatment effects
Target: $\hat{y}(x) = E[Y_i   X_i = x]$	Target: $\tau(x) = E[\tau_i   X_i = x]$
Criterion:	Criterion:
$\min E[(\hat{y}(x) - Y_i)^2   X_i = x]$	$\min E[(\tau(x) - \tau_i)^2   X_i = x]$
Training data: $\{Y_i, X_i\}_{i=1}^n$	Training data: $\{\tau_i, X_i\}_{i=1}^n$

Why is training data a problem for predicting treatment effects?

- ▶ Consequence: can't apply ML directly to predicting treatment effects; have to adapt them

## Adapting ML to predict treatment effects

- ▶ Break it up:

$$\begin{aligned} E[\tau_i | X_i] &:= E[Y_i(1) - Y_i(0) | X_i] \\ &= E[Y_i | X_i, D_i = 1] - E[Y_i | X_i, D_i = 0] \end{aligned}$$

(by what assumption?)

## Adapting ML to predict treatment effects

- ▶ Break it up:

$$\begin{aligned} E[\tau_i | X_i] &:= E[Y_i(1) - Y_i(0) | X_i] \\ &= E[Y_i | X_i, D_i = 1] - E[Y_i | X_i, D_i = 0] \end{aligned}$$

(by what assumption?)

- ▶ Adjust the criterion: (why?)

$$\min \sum_{i=1}^n (\tau(X_i) - \tau_i)^2 \iff \max \sum_{i=1}^n \tau(X_i)^2$$

# Adapting ML to predict treatment effects

- ▶ Break it up:

$$\begin{aligned} E[\tau_i | X_i] &:= E[Y_i(1) - Y_i(0) | X_i] \\ &= E[Y_i | X_i, D_i = 1] - E[Y_i | X_i, D_i = 0] \end{aligned}$$

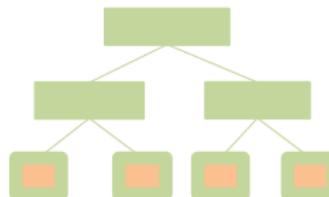
(by what assumption?)

- ▶ Adjust the criterion: (why?)

$$\min \sum_{i=1}^n (\tau(X_i) - \tau_i)^2 \iff \max \sum_{i=1}^n \tau(X_i)^2$$

- ▶ Be honest: use one set of observations to select the tree structure, and another to generate predictions

Y	x1	x2	x3



# Predicting treatment effects using ML: Summary

- ▶ Target:

$$CATE := \tau(x) = E[\tau_i | X_i = x]$$

- ▶ Key identifying assumption:

$$(Y_i(0), Y_i(1)) \perp\!\!\!\perp D_i | X_i$$

- ▶ Estimation: Random Causal Forest

- ▶ Grow decision trees on many bootstrapped samples
- ▶ Choose splits using the training set to max  $\sum_{i=1}^n \tau(X_i)^2$
- ▶ Generate predictions in each leaf using the estimation set
- ▶ Average predictions over the trees in the forest

- ▶ Go to python!

# Thank you!